

DPDP RAG — KPI Report

Citation-grounded DPDP content system · Report date 21 Jul 2026 · Eval: 18-case golden set (default-prose path, judge = gpt-4o-mini) · Operational: LangSmith default, last 30 days

Overall: GREEN. Every non-negotiable golden-rule gate passes at 100%. No fabricated citations, honest effective-dating, strong retrieval and faithfulness. The only recorded errors are the resolved 15–16 Jul OpenAI-key incident; zero errors since 17 Jul. The system is conservative by design — it escalates coverage gaps to human review rather than auto-passing.

1 · Golden-rule gates — non-negotiable

GATE	VALUE	TARGET	STATUS
citation_validity (no fabricated citations)	100%	100%	PASS
date_honesty (future provisions dated)	100%	100%	PASS
no_contradiction	100%	100%	PASS
gate_refusal (junk queries killed pre-LLM)	100%	100%	PASS
no_false_assertion	100%	100%	PASS
answered	100%	100%	PASS

2 · Retrieval quality (k = 6)

METRIC	VALUE	TARGET
recall@k	100%	80%
MRR	88.0%	70%
nDCG@k	88.3%	75%
HitRate	100%	—
R-precision	66.7%	ungated

3 · Generation quality (LLM-judged)

METRIC	VALUE	TARGET
faithfulness (≤ sources)	98.9	85
correctness	100	85
relevance	100	85
coherence	90.6	80
conciseness	84.2	75

4 · Traceability & self-verification

Per-draft traceability: mean 76.4% (range 54–100%) — the share of factual claims carrying a citation. **11 of 18** drafts fell below the 90% bar and **escalated to human review**; machine auto-pass rate (verified) = **38.9%**.

This is **by design, not a defect**. citation_validity is 100% — *nothing is fabricated*. The gap is *uncited* claims (largely rhetorical framing), and the pipeline **refuses to auto-pass them — it routes to a human reviewer**. Escalation is the safety net doing its job. Styled LinkedIn formats trace similarly (54–77%) and also escalate.

5 · Operational — LangSmith, last 30 days

METRIC	VALUE	NOTE
Root runs	1,031	1,007 local dev/eval · 24 prod (Render)
Error rate	1.6%	all errors 15–16 Jul (key incident); 0 since 17 Jul
Prod latency (Render)	P50 21.2s	P95 57.5s
Eval latency	gen P50 16.7s	retrieval P50 3.3s · gen P95 23.1s

LLM calls	785	OpenAI 454 · Anthropic 331 (42%)
Tracing	Active	.env + Render prod both on

6 · Watch-items

- **Conservative escalation (~61% of drafts).** Coverage below 90% routes to human review. Correct for legal content; raising the auto-pass rate is a policy dial (drafter citing more of its framing sentences), not a correctness fix.
- **Anthropic = 42% of LLM calls over 30 days.** Likely failover during the 15–16 Jul OpenAI outage plus the retrieval workflow's Claude decompose/sufficiency calls. Worth confirming *recent* failover is low now that OpenAI is stable.
- **R-precision 66.7%** (ungated): ordering inside the top-6 — all 6 chunks still reach the prompt, so it does not affect grounding.

Sources: project eval harness (`eval/run_eval.py` → `eval/last_run.json`, RAGAS-style dimensions computed directly + golden-rule gates) and LangSmith run analytics. Golden-rule thresholds per project spec (faithfulness ≥ 85, recall@k ≥ 80). AI-assisted report; figures reflect the last measured eval run and a 30-day operational window.